

---

## HPC for numerical methods and data analysis

*Fall Semester 2023*

*Prof. Laura Grigori*

*Assistant: Mariana Martínez Aguilar*

**Session 2 – September 16, 2024**

---

## Clusters, Python, and MPI

### 1 Clusters

High performance computing (HPC) is the ability to process (usually huge amounts of) data and perform complex calculations at high speed. This is important in today's world because of the increase in available data. To be able to do this, HPC makes use of clusters of powerful processors that work in parallel. These systems typically run at much higher speeds than commercially available laptops.

A server is a computer that provides information to other computers called "clients" on computer networks. Clusters are groups of servers that are managed together and participate in workload management. A cluster can contain nodes or individual application servers. This depends on the type of cluster. Clusters are responsible for balancing workload among servers. Servers that are a part of a cluster are called cluster members. When an application is installed on a cluster, it is automatically installed on each cluster member. This is why we can distribute client tasks in distributed platforms according to the capabilities of the different machines by assigning weights to each server. In distributed platforms, assigning weights to the servers in a cluster improves performance and failover. Tasks are assigned to servers that have the capacity to perform those task operations but if one server is unable to perform the task, it can be reassigned.

A node is a computer part of a large set of nodes (cluster). A computer node offers different types of resources: processors, volatile memory (RAM), permanent disc space (SSD), accelerators (GPUs), etc. A node group defines groups of nodes that are capable of hosting members of the same cluster. By organising nodes that satisfy an application requirements into a node group, we establish an administrative policy that governs which nodes can be used together to form a cluster. Nodes can be members of multiple groups.

A core is the part of a processor that does the computation. A processor comprises multiple cores, as well as a memory controller, a bus controller, and other components. A core group is a group of clusters in a high availability environment. All of the application servers defined as a member of one of the clusters included in a core group are automatically members of that core group.

Other than the applications configured to run on them, cluster members do not have to share any other configuration data. This allows client work to be distributed across all the members of a cluster instead of all workload being handled by a single application server. A vertical cluster has cluster members on the same physical machine while a horizontal cluster has cluster members on

multiple nodes across many machines.

A workload manager like Slurm is designed to provide the system administrator with increased control over how the scheduler virtual memory manager (VMM) and the disc I/O subsystem allocate resources to processes.

We're going to make sure that we can correctly connect to the cluster available for this course. If you have more questions you can read the documents in depth here. **Make sure you have a GASPARE account**, let your username be `jusernamej`. The cluster available for us is called helvetios. Take into consideration that you can only connect if you're physically at EPFL, otherwise you'll need to use a VPN.

### Exercise I Reminder of a simple MPI code in Python

Given two vectors,  $b, c$  we want to compute  $d = 2b + c$ . Execute the following simple code on 2 processors several times.

```
from mpi4py import MPI
import numpy as np

b = np.array([1, 2, 3, 4])
c = np.array([5, 6, 7, 8])
a = np.zeros_like(b)
d = np.zeros_like(b)

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    for i in range(4):
        a[i] = b[i] + c[i]
    comm.Send(a, dest = 1, tag = 77)
else:
    comm.Recv(a, source = 0, tag = 77)
    for i in range(4):
        d[i] = a[i] + b[i]

print("I am rank = ", rank )
print("d: ", d)
```

Observe the order in which the prints take place and the value of  $d$  at the end.

### Exercise II Point to point communication - blocking and non-blocking communication

- a) Provide a brief definition of MPI. What is a communicator?
- b) Execute the following simple code on 4 processors.

```
from mpi4py import MPI
```

```

import numpy as np

# Initialize MPI
comm = MPI.COMM_WORLD
rank = comm.Get_rank()

if rank == 0:
    data = {'a': 7, 'b': 3.14}
    print("From process: ", rank, "\n data sent:", data, "\n")
    comm.send(data, dest=1, tag=11)
elif rank == 1:
    data = comm.recv(source=0, tag=11)
    print("From process: ", rank, "\n data received:", data, "\n")
elif rank == 2:
    data = np.array([1, 1, 1, 1, 1])
    print("From process: ", rank, "\n data sent:", data, "\n")
    comm.send(data, dest=3, tag = 66)
else:
    data = comm.recv(source = 2, tag = 66)
    print("From process: ", rank, "\n data received:", data, "\n")

```

In this case, why do we need to be careful when specifying the `dest` and `tag` parameters on both `comm.send` and `comm.recv`?

- c) Describe the difference between blocking communication and non-blocking communication in MPI. Modify the code above such that it uses `comm.isend` instead of `comm.send` and `comm.irecv` instead of `comm.recv` while ensuring the messages are passed correctly.

### Exercise III Collective communication - scattering and broadcasting

- a) Run the following script on 4 processors:

```

from mpi4py import MPI
import numpy as np

# Initialize MPI
comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

# Define the vector
if rank == 0:
    vector = np.array([16, 62, 97, 25])
else:
    vector = None

```

```

data1 = comm.bcast(vector, root = 0)
data2 = comm.scatter(vector, root = 0)

print("rank: ", rank, " data1: ", data1, " data2: ", data2)

```

What is the difference in MPI between scattering and broadcasting?

b) Consider the multiplication of a matrix  $A \in \mathbb{R}^{m \times n}$  with a vector  $v \in \mathbb{R}^n$ . Write a Python file containing a script that:

- Creates a matrix of dimension  $m \times n$
- Creates a vector of dimension  $n$
- Makes sure that the dimensions of the matrix and the vector agree in such way that we can compute  $Av$
- Computes  $Av$  using MPI's scattering, make sure you execute your code on the right amount of processors (*Hints: you'll need to use `comm.gather`. What are the entries of  $Av$ ?*)

#### Exercise IV Collective communication - all-to-all and reduce

- Run the following code on 4 processors:

```

from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
size = comm.Get_size()
rank = comm.Get_rank()

senddata = rank*np.ones(size, dtype = int)

recvdata = comm.alltoall(senddata)

print(" process ", rank, " sending ", senddata, " receiving ", recvdata )

```

What is `comm.alltoall` doing? Compare it to `comm.scatter`.

- In this exercise we are going to use reduction operations on MPI. Run the following code on 4 processors:

```

from mpi4py import MPI
import numpy as np

comm = MPI.COMM_WORLD
size = comm.Get_size()
rank = comm.Get_rank()

```

```

senddata = rank*np.ones(size, dtype = int)

global_result1 = comm.reduce(senddata, op = MPI.SUM, root = 0)
global_result2 = comm.reduce(rank, op = MPI.MAX, root = 0)

#Print
print(" process ", rank, " sending ", senddata)

#Print the result on the root process
if rank == 0:
    print(" Reduction operation1: ", global_result1,
          "\n Reduction operation2: ", global_result2)

```

What is a reduction operation? What is the difference between this and `comm.gather`?

- In the previous code, change `comm.reduce` to `comm.allreduce`. What is the difference between the two? (Note, `comm.allreduce` doesn't use the argument `root`).

## Exercise V Deciding what to use - Mid point rule

Numerical integration describes a family of algorithms for calculating the value of definite integrals. One of the simplest algorithms to do so is called the Mid Point Rule. Assume that  $f(x)$  is continuous on  $[a, b]$ . Let  $n$  be a positive integer and  $h = (b - a)/n$ . If  $[a, b]$  is divided into  $n$  subintervals,  $\{x_0, x_1, \dots, x_{n-1}\}$ , then if  $m_i = (x_i + x_{i+1})/2$  is the midpoint of the  $i$ -th subinterval, set:

$$M_n = \sum_{i=1}^n f(m_i)h.$$

Then:

$$\lim_{n \rightarrow \infty} M_n = \int_a^b f(x)dx.$$

Thus, for a fixed  $n$ , we can approximate this integral as:

$$\int_a^b f(x)dx \approx \sum_{i=1}^n f(m_i)h$$

Set  $n = s * 500$ ,  $f(x) = \cos(x)$ ,  $a = 0$ ,  $b = \pi/2$ . Write a Python script such that:

- Defines a function that given  $x_i, h, n$  first calculates 500 mid points on a subinterval  $[x_i, x_{i+1}]$  and returns the approximation of the integral on this subinterval.
- Using MPI approximates the integral of  $f$  on  $[a, b]$
- Run your script on  $s$  processors

to approximate the integral of  $f$ . (*Hints: there are many ways of doing this, one approach is using `comm.bcast` and `comm.reduce`*).