
HPC for numerical methods and data analysis

Fall Semester 2023

Prof. Laura Grigori

Assistant: Mariana Martinez

Session 7 – October 31, 2023

Randomized SVD

Exercise 1: SRHT

In the context of overdetermined least-squares problems, we need to find $x \in \mathbb{R}^n$ such that it minimizes:

$$\|Ax - b\|_2^2,$$

where $A \in \mathbb{R}^{m \times n}, b \in \mathbb{R}^m, m > n$. There is a class of randomized algorithms for solving this problem based on sketching method. Sketching methods involve using a random matrix $\Omega \in \mathbb{R}^{r \times m}$ to project the data A (and maybe also b) to a lower dimensional space with $r \ll m$. Then they approximately solve the least-squares problem using the sketch ΩA (and/or Ωb). One relaxes the problem to finding a vector x so that

$$\|Ax - b\| \leq (1 + \varepsilon)\|\Omega A x - \Omega b\|,$$

where x^* is the optimal solution. The overview of sketching applied to solve linear least squares is:

- Sample/build a random matrix Ω
- Compute ΩA and Ωb
- Output the exact solution to the problem $\min_x \|(\Omega A)x - (\Omega b)\|_2$.

Given a data matrix, $X \in \mathbb{R}^{m \times n}$, we want to reduce the dimensionality of X by defining a random orthonormal matrix $\Omega \in \mathbb{R}^{r \times m}$ with $r \ll m$. For $m = 2^q, q \in \mathbb{N}$, the Subsampled Randomized Hadamard Transform (SRHT) algorithm defined a $r \times m$ matrix as:

$$\Omega = \sqrt{\frac{m}{r}} P H_m D,$$

where:

- $D \in \mathbb{R}^{m \times m}$ is a diagonal matrix whose elements are independent random signs, i.e. its diagonal entries are just -1 or 1 .

- $H \in \mathbb{R}^{m \times m}$ is a **normalized** Walsh-Hadamard matrix. If you're going to use a library that implements this transform then check that it implements the normalized Walsh-Hadamard matrix. This matrix is defined recursively as:

$$H_m = \begin{bmatrix} H_{m/2} & H_{m/2} \\ H_{m/2} & -H_{m/2} \end{bmatrix} \quad H_2 = \begin{bmatrix} 1 & 1 \\ 1 & -1 \end{bmatrix}$$

$$H = \frac{1}{\sqrt{m}} H_m \in \mathbb{R}^{m \times m}.$$

- $P \in \mathbb{R}^{r \times m}$ is a subset of randomly sampled r columns from the $m \times m$ identity matrix. The purpose of using P is to uniformly sample r columns from the rotated data matrix $X_{\text{rot}} = H_m D X$.

The following theorem help us get an idea for the size of r .

Theorem 1 (Subsampled Randomized Hadamard Transform) Let $\Omega = \sqrt{\frac{m}{r}} P H_m D$ as previously defined. Then if

$$r \geq \mathcal{O}((\varepsilon^{-2} \log(n))(\sqrt{n} + \sqrt{\log m})^2)$$

with probability 0,99 for any fixed $U \in \mathbb{R}^{m \times n}$ with orthonormal columns:

$$\|I - U^\top \Omega \Omega^\top U\|_2 \leq \varepsilon.$$

Further, for any vector $x \in \mathbb{R}^m$, Ωx can be computed in $\mathcal{O}(n \log r)$ time.

Choose a data set from [<https://www.kaggle.com/datasets/tags=13405-Linear+Regression>]. Compare the randomized least squares fit using SRHT vs the deterministic least squares fit. Use the previous theorem to estimate r . *Hint: you can use the fast Hadamard transform from scipy or pytorch*

Solution: Below you can find an implementation of SRHT:

```
import numpy as np
from numpy.linalg import norm, lstsq
from pandas import read_csv
from numpy.random import normal
from math import ceil, log, sqrt, floor
import matplotlib.pyplot as plt
import time
from random import sample
import random
import torch
from hadamard.transform import hadamard_transform

plt.ion()

# For SRHT sketching applied to a least squares problem
# we report the following quantities:
##### Time taken to solve the full problem
##### Time taken to solve the compressed problem
##### Residual norm full problem
##### Residual norm compressed problem
##### Relative error in the spectral norm
```

```

# We are going to read the data (which was previously downloaded)
# We just want to work with certain columns, not all of them
d = read_csv("ParisHousing.csv")
b = d.price
b = b.values
d.drop(['hasYard', 'hasPool', 'floors', 'cityCode', 'numPrevOwners',
        'made', 'basement', 'attic', 'garage', 'hasGuestRoom'], axis = 1)
A = d.values
# But we need to make sure m is a power of 2
m = int(2**(floor(log(A.shape[0])/log(2))))
A = A[0:m, :]
b = b[0:m]

# Now that we have out set up
n = A.shape[1]
nRuns = 10
sigma = 0.99
epsilon = np.array([100, 10, 5, 2, 1, 0.5, 0.1])
rVec = np.ceil((log(n)/(epsilon**2))*(sqrt(n) + log(m)**2).astype('int'))
# Notice that some r's might be bigger than m

timeF = np.empty_like(epsilon)
timeC = np.empty_like(epsilon)
resF = np.empty_like(epsilon)
resC = np.empty_like(epsilon)
relErrSpec = np.empty_like(epsilon)

for k in range(len(epsilon)):
    eps = epsilon[k]
    r = min(m, rVec[k])
    tF = 0
    tC = 0
    rC = 0
    rES = 0
    for run in range(nRuns):
        # Begin with the compressed problem
        ts = time.time()
        d = np.array([1 if random.random() < 0.5 else -1 for i in range(m)])
        D = np.diag(sqrt(m/r)*d)
        P = sample(range(m), r)
        omega = D
        omega = np.array([hadamardtransform(torch.from_numpy(omega[:, i])).numpy() for i in range(m)])
        omega = np.transpose(omega)
        omega = omega[P, :]
        omegaA = omega@A
        omegab = omega@b
        xPrime = lstsq(omegaA, omegab)
        xPrime = xPrime[0]
        tC += time.time() - ts
        # Now for the full problem
        ts = time.time()
        xStar = lstsq(A, b)
        xStar = xStar[0]
        tF += time.time() - ts
        # Report desired quantities for the randomized part
        rC += norm(omegaA@xPrime - omegab)
        rES += abs(norm(omegaA) - norm(A))/norm(A)

```

```

# Save averages
timeF[k] = tF/nRuns
timeC[k] = tC/nRuns
resF[k] = norm(A@xStar - b)
resC[k] = rC/nRuns
relErrSpec[k] = rES/nRuns

###
### Plot plot plot
# Time
plt.figure(figsize=(8, 6), dpi=80)
plt.loglog(epsilon, timeF, c = "#003aff", marker = 'o',
           label = "Full problem")
plt.loglog(epsilon, timeC, c = "#00b310", marker = '*',
           label = "Compressed problem")
plt.legend()
plt.title(r'$\varepsilon$' +
          ", time taken to build and compute")
plt.xlabel(r'$\varepsilon$')
plt.ylabel("Time, s")

# Norm of residual
plt.figure(figsize=(8, 6), dpi=80)
plt.loglog(epsilon, resF, c = "#003aff", marker = 'o',
           label = "Full problem")
plt.loglog(epsilon, resC, c = "#00b310", marker = '*',
           label = "Compressed problem")
plt.legend()
plt.title(r'$\varepsilon$' + ", norm of residual")
plt.xlabel(r'$\varepsilon$')
plt.ylabel("Norm of residual")

# Relative error in spectral norm
plt.figure(figsize=(8, 6), dpi=80)
plt.loglog(epsilon, relErrSpec, c = "#5400b3", marker = 'o',
           label = "Relative error")
plt.loglog(epsilon, epsilon, c = '#676b74', linestyle='dashed',
           label = r'$\varepsilon$')
plt.legend()
plt.title(r'$\varepsilon$' + ", relative error spectral norm " +
          r'$\|\Omega A\|_2 - \|A\|_2 / \|A\|_2$')
plt.xlabel(r'$\varepsilon$')
plt.ylabel(r'$\|\Omega A\|_2 - \|A\|_2 / \|A\|_2$')

```

Exercise 2: Randomized SVD

Rokhlin, Szlam, and Tygert introduced an algorithm called *Blanczos* such that it computes the whole approximation $U\Sigma V^\top$ to an SVD of a matrix $A \in \mathbb{R}^{m \times n}$.

Test this algorithm by constructing a rank- k approximation with $k = 10$ to a matrix $A \in \mathbb{R}^{m \times 2m}$ via its SVD:

$$A = U^{(A)} \Sigma^{(A)} V^{(A)\top},$$

where:

- $U \in \mathbb{R}^{m \times m}$ is a Hadamard matrix

Algorithm 1 Blanczos

Input: $A \in \mathbb{R}^{m \times n}$, i, l such that $k < l$ and $(i+1)l \leq m - k$

Output: U, Σ, V

Form a real $l \times n$ matrix G such that its entries are i.i.d. Gaussian random variables with mean zero and unit variance. Compute:

$$\begin{aligned} R^{(0)} &= GA \\ R^{(1)} &= R^{(0)} A^\top A \\ &\vdots \\ R^{(i)} &= R^{(i-1)} A^\top A. \end{aligned}$$

Form the $(i+1)l \times n$ matrix:

$$R^\top = [(R^{(0)})^\top \quad (R^{(1)})^\top \quad \dots \quad (R^{(i)})^\top]$$

Form a real $n \times (i+1)l$ matrix Q whose columns are orthonormal and such that there is a real $(i+1)l \times (i+1)l$ matrix S in such way that $R^\top = QS$

$T \leftarrow AQ$

Form the SVD of T , $T = U\Sigma W^\top$

$V \leftarrow QW$

- $V \in \mathbb{R}^{2m \times 2m}$ is a Hadamard matrix
- $\Sigma \in \mathbb{R}^{m \times 2m}$ is a diagonal matrix whose diagonal entries are defined as:

$$\Sigma_{jj} = \sigma_j = (\sigma_{k+1})^{\lfloor j/2 \rfloor / 5},$$

for $j = 1, 2, \dots, 9, 10$ and

$$\Sigma_{jj} = \sigma_j = \sigma_{k+1} \frac{m-j}{m-11},$$

for $j = 11, 12, \dots, m-1, m$. Thus $\sigma_1 = 1$ and $\sigma_k = \sigma_{k+1}$.

Set $l = k+12, i = 1$ test this algorithm for $m = 2^{11}$, $\sigma_{k+1} = 0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001$. Plot the decay of the singular values of A and compare such decay with the accuracy of the approximation, $\|A - U\Sigma V^\top\|_F$ and the relative error, $\frac{\|A - U\Sigma V^\top\|_F}{\|A\|_F}$.

Solution: Below you can find a script with implementation of Blanczos and generation of necessary plots:

```
import numpy as np
from numpy.linalg import svd, qr, norm
import matplotlib.pyplot as plt
from scipy.linalg import hadamard
from math import log, sqrt, floor
import torch
from hadamard.transform import hadamard.transform

# So that the plots are "interactive" when we run this script
plt.ion()
```

```

def SVD_Blanczos(A, l, i, Sigma_full = False):
    '''
    From Rokhlin, Szlam, Tygert paper A Randomized Algorithm For Principal Component Analysis, algorithm 1
    IN :
        A          : mxn matrix to be factorized
        l          : paramter for our approximated matrix C of size mxc
        i          : order of approximation wanted
        Sigma_full  : transition probabilities
    OUT :
        U          : approximated left singular vectors
        Sigma       : approximated singular values
        V          : approximated right singular vectors
    '''
    m = A.shape[0]
    n = A.shape[1]
    # STEP 1
    # Using a random number generator form a real lxm matrix G whose entries are iid Gaussian and complex
    G = np.random.normal(loc= 0.0, scale = 1.0, size = [l, m])
    R = np.zeros(( (i+1)*l, n ))
    R_temp = G@A
    R[0:(l), :] = R_temp
    for j in range(i):
        R_temp = R_temp@np.transpose(A)@A
        R[ (j+1)*l:(j+2)*l, : ] = R_temp
    # STEP 2
    # Using QR decomposition form a real n x ((i+1)l) matrix Q whose columns are orthonormal
    Q, S = qr(np.transpose(R))
    # STEP 3
    # Compute the m x ( (i+1)l ) product matrix
    T = A@Q
    # STEP 4
    # Form an SVD of T
    U, Sigma, Wt = svd(T)
    # STEP 5
    # Compute the n x ( (i+1)l ) product matrix
    V = Q@np.transpose(Wt)
    if Sigma_full:
        S_t = Sigma
        Sigma = np.zeros((m,n))
        np.fill_diagonal(Sigma, S_t)
    return U, Sigma, V

def buildA(m, sigma_kl, k = 10):
    '''
    From Rokhlin, Szlam, Tygert paper A Randomized Algorithm For Principal Component
    Analysis, build test matrix A of size mx(2m). We use the fast Hadamard transform
    IN:  m          : number of desired rows in matrix A
        sigma_kl    : (k+1)th biggest singular value of A
        k          : where we are going to truncate the approximation of A
    OUT: A          : matrix with desired structure

    QUESTION: Can we build A faster? Notice that Sigma is just a diagonal matrix.
    Also notice that we can use the fast Hadamard transform to build A.
    If you can, change this function so that it builds A faster!
    '''
    U = (1/sqrt(m)) * hadamard(m)
    V = (1/sqrt(2*m)) * hadamard(2*m)

```

```

firstSig = [sigma_k1**((floor(j/2)/5) for j in range(1, k+1)]
sigmas = firstSig + [sigma_k1*(m - j)/(m - 11) for j in range(k+1, m+1)]
Sigma = np.zeros((m, 2*m))
np.fill_diagonal(Sigma, sigmas)
return U@Sigma@np.transpose(V), sigmas

# Test
m = 2**11
k = 10
l = k + 12
i = 1
sigma_k1S = [0.1, 0.01, 0.001, 0.0001, 0.00001, 0.000001]
errorApprox = np.empty(6)
errorApproxRel = np.empty(6)

for s in range(6):
    sigma = sigma_k1S[s]
    # Build A
    A, sigmas = buildA(m, sigma, k)
    # Blenczos
    U, S, V = SVD.Blenczos(A, l, i)
    Sigma = np.zeros((U.shape[1], S.shape[0]))
    np.fill_diagonal(Sigma, S)
    # Plot the decay of the singular values
    plt.figure(figsize=(8, 6), dpi=80)
    plt.loglog(np.arange(m), sigmas, marker = 'o', c = "#0800ff")
    plt.title("Decay on singular values for " + r"$\sigma_{k+1} = $" + str(sigma))
    plt.xlabel("k")
    plt.ylabel(r"$\sigma_{k}$")
    # Save the error of the approximation || A - U \Sigma V^top ||
    errorApprox[s] = norm(A - U@Sigma@np.transpose(V), 'fro')
    errorApproxRel[s] = errorApprox[s]/norm(A, 'fro')

# Plot the errors of the approximation
plt.figure(figsize=(8, 6), dpi=80)
plt.loglog(sigma_k1S, errorApprox, marker = 'o', c = "#8700ff")
plt.title("Errors in approximation, different " + r"$\sigma_{k+1}$")
plt.xlabel(r"$\sigma_{k+1}$")
plt.ylabel(r"$\| A - U \Sigma V^top \| $")

# Plot relative errors of the approximation
plt.figure(figsize=(8, 6), dpi=80)
plt.loglog(sigma_k1S, errorApproxRel, marker = 'o', c = "#ff8f00")
plt.title("Relative errors in approximation, different " + r"$\sigma_{k+1}$")
plt.xlabel(r"$\sigma_{k+1}$")
plt.ylabel(r"$\| A - U \Sigma V^top \| $")

```